

Slub_leak_debug

Some platforms maybe too old or too new, we cannot use slab trace.

Or sometimes the page_owner is not available. Here I get a method to debug slub memleak.

1: first you need to confirm the slub is keep leaking.

Cat /proc/meminfo

Then check the SUnreclaim for exp:

Here is we cat /proc/meminfo every 2-3 hours, then we found the SUnreclaim is keep growing

行 26: SUnreclaim:	781988 kB
行 71: SUnreclaim:	806644 kB
行 116: SUnreclaim:	815900 kB
行 161: SUnreclaim:	829012 kB
行 206: SUnreclaim:	838196 kB
行 251: SUnreclaim:	851020 kB
行 296: SUnreclaim:	873748 kB
行 341: SUnreclaim:	873476 kB
行 386: SUnreclaim:	889788 kB
行 431: SUnreclaim:	904212 kB
行 476: SUnreclaim:	906100 kB
行 520: SUnreclaim:	907468 kB
行 566: SUnreclaim:	917092 kB
行 611: SUnreclaim:	947828 kB
行 656: SUnreclaim:	949708 kB
行 701: SUnreclaim:	951100 kB
行 746: SUnreclaim:	954340 kB
行 791: SUnreclaim:	955300 kB
行 836: SUnreclaim:	956996 kB
行 881: SUnreclaim:	958532 kB
行 926: SUnreclaim:	960140 kB
行 971: SUnreclaim:	962076 kB
行 1016: SUnreclaim:	968716 kB
行 1061: SUnreclaim:	968276 kB
行 1106: SUnreclaim:	973860 kB
行 1151: SUnreclaim:	977180 kB
行 1196: SUnreclaim:	982332 kB
行 1241: SUnreclaim:	986188 kB
行 1286: SUnreclaim:	989196 kB
行 1331: SUnreclaim:	990236 kB

```

行 1376: SUnreclaim:      992356 kB
行 1421: SUnreclaim:      994628 kB
行 1466: SUnreclaim:      994820 kB
行 1510: SUnreclaim:     1006092 kB

```

2: meanwhile we need also check the slab info in /proc/slabinfo every 2-3 hours,

Then we can find below slab is keep growing

anon_vma	235674	235725	456	35	4 : tunables	0	0	0 : slabdata	6735	6735	0
anon_vma	256207	256235	456	35	4 : tunables	0	0	0 : slabdata	7321	7321	0
anon_vma	267890	267890	456	35	4 : tunables	0	0	0 : slabdata	7654	7654	0
anon_vma	275587	275590	456	35	4 : tunables	0	0	0 : slabdata	7874	7874	0
anon_vma	279056	279090	456	35	4 : tunables	0	0	0 : slabdata	7974	7974	0
anon_vma	289809	289835	456	35	4 : tunables	0	0	0 : slabdata	8281	8281	0
anon_vma	307633	307650	456	35	4 : tunables	0	0	0 : slabdata	8790	8790	0
anon_vma	311486	311465	456	35	4 : tunables	0	0	0 : slabdata	8899	8899	0
anon_vma	324169	324170	456	35	4 : tunables	0	0	0 : slabdata	9262	9262	0
anon_vma	352474	352485	456	35	4 : tunables	0	0	0 : slabdata	10071	10071	0
anon_vma	355662	355670	456	35	4 : tunables	0	0	0 : slabdata	10162	10162	0
anon_vma	358776	358785	456	35	4 : tunables	0	0	0 : slabdata	10251	10251	0
anon_vma	372857	372890	456	35	4 : tunables	0	0	0 : slabdata	10654	10654	0
anon_vma	405199	405230	456	35	4 : tunables	0	0	0 : slabdata	11578	11578	0
anon_vma	409992	410025	456	35	4 : tunables	0	0	0 : slabdata	11715	11715	0
anon_vma	413193	413245	456	35	4 : tunables	0	0	0 : slabdata	11807	11807	0
anon_vma	423251	423290	456	35	4 : tunables	0	0	0 : slabdata	12094	12094	0
anon_vma	426559	426580	456	35	4 : tunables	0	0	0 : slabdata	12188	12188	0
anon_vma	430961	430990	456	35	4 : tunables	0	0	0 : slabdata	12314	12314	0
anon_vma	436349	436380	456	35	4 : tunables	0	0	0 : slabdata	12468	12468	0
anon_vma	439311	439320	456	35	4 : tunables	0	0	0 : slabdata	12552	12552	0
anon_vma	442956	442995	456	35	4 : tunables	0	0	0 : slabdata	12657	12657	0
anon_vma	450505	450520	456	35	4 : tunables	0	0	0 : slabdata	12872	12872	0
anon_vma	453395	453425	456	35	4 : tunables	0	0	0 : slabdata	12955	12955	0
anon_vma	466632	466690	456	35	4 : tunables	0	0	0 : slabdata	13334	13334	0
anon_vma	471149	471170	456	35	4 : tunables	0	0	0 : slabdata	13462	13462	0
anon_vma	480762	480795	456	35	4 : tunables	0	0	0 : slabdata	13737	13737	0
anon_vma	487191	487235	456	35	4 : tunables	0	0	0 : slabdata	13921	13921	0
anon_vma	491550	491610	456	35	4 : tunables	0	0	0 : slabdata	14046	14046	0
anon_vma	496472	496475	456	35	4 : tunables	0	0	0 : slabdata	14185	14185	0
anon_vma	499981	500010	456	35	4 : tunables	0	0	0 : slabdata	14286	14286	0
anon_vma	503502	503510	456	35	4 : tunables	0	0	0 : slabdata	14386	14386	0

3: next we can add trace on the slab alloc function.

```
diff --git a/mm/slub.c b/mm/slub.c
```

```
index 229136a8ef70..cb05adc1ca54 100644
```

```
--- a/mm/slub.c
```

```
+++ b/mm/slub.c
```

```
@@ -229,6 +229,27 @@ static inline void stat(const struct kmem_cache *s, enum stat_item si)
```

```
#endif
```

```
}
```

```
+void trace_my_debug_slub_alloc(struct kmem_cache *s)
```

```
+{
```

```
+if((0 == strcmp("anon_vma",s->name))){ // change this to trace your leak slab
+ trace_printk(
+ "alloc slub %s = Callers:(%ps<-%ps<-%ps<-%ps<-%ps<-%ps)\n",s->name,
+ (void *)CALLER_ADDR0, (void *)CALLER_ADDR1, (void *) CALLER_ADDR2,
+ (void *)CALLER_ADDR3, (void *)CALLER_ADDR4, (void *)CALLER_ADDR5,
+ (void *)CALLER_ADDR6);
+}
+}
+
+void trace_my_debug_slub_free(struct kmem_cache *s)
+{
+if((0 == strcmp("anon_vma",s->name))){
+ trace_printk(
+ "free slub %s = Callers:(%ps<-%ps<-%ps<-%ps<-%ps<-%ps)\n",s->name,
+ (void *)CALLER_ADDR0, (void *)CALLER_ADDR1, (void *) CALLER_ADDR2,
+ (void *)CALLER_ADDR3, (void *)CALLER_ADDR4, (void *)CALLER_ADDR5,
+ (void *)CALLER_ADDR6);
+}
+}
+}
+
+*****
+
+* Core slab cache functions
+***** /
@@ -2969,6 +2990,7 @@ void *kmem_cache_alloc(struct kmem_cache *s, gfp_t gfpflags)
trace_kmem_cache_alloc(_RET_IP_, ret, s->object_size,
s->size, gfpflags);

+ trace_my_debug_slub_alloc(s);

return ret;
}

EXPORT_SYMBOL(kmem_cache_alloc);

@@ -3225,6 +3247,7 @@ void kmem_cache_free(struct kmem_cache *s, void *x)

return;

slab_free(s, virt_to_head_page(x), x, NULL, 1, _RET_IP_);

trace_kmem_cache_free(_RET_IP_, x);

+ trace_my_debug_slub_free(s);

}

EXPORT_SYMBOL(kmem_cache_free);
```

```
echo 0 > /sys/kernel/debug/tracing/tracing_on
```

```
cat /sys/kernel/debug/tracing/trace_pipe > /sdcard/fttrace.txt & // get trace log
```

[illegible]

```
./sort_slab ftrace_15.txt // 最新版本已做数据排序
```

```
-----
---total alloc 1457855-----
---total free 88-----
---total leak 1457767-----
```

```

/*
* 使用方法，编译好后 如 gcc sort2.c -o debug_slab_sort
*
* ./debug_slab_sort ftrace.txt
*
* 其中可自行调整 printList 函数来决定输出的数据， 目前已做数据排序
*
*/

```

```

#include <stdio.h>
#include <string.h>
#include <stdlib.h>
typedef struct Node {
    char name[100];
    unsigned long int alloc_count;
    unsigned long int free_count;
    struct Node* next;
} Node;
#define LEAK_TIMES 0
Node* head = NULL;
Node* phead = NULL;
int check_alloc_or_free(char * strings) {
    char str1[] = "trace_my_debug_slub_alloc";
    char str2[] = "trace_my_debug_slub_free";
    char *ptr = strstr(strings, str1);
    if (ptr) {
        return 1;
    }
    ptr = strstr(strings, str2);
    if (ptr) {
        return 2;
    }
    return 0;
}
char* get_task(char * strings) {
    const char s[2] = "[";

```

```

    char *token;
    token = strtok(strings, s);
    return token;
}
Node* createNode(char* name) {
    Node* newNode = (Node*)malloc(sizeof(Node));
    if (newNode == NULL) {
        printf("Error! Unable to allocate memory.\n");
        return NULL;
    }
    strcpy(newNode->name, name);
    newNode->alloc_count = 0;
    newNode->free_count = 0;
    newNode->next = NULL;
    return newNode;
}
void appendNode(char* name, int state) {
    Node* newNode = createNode(name);
    Node* temp = NULL;
    if (head == NULL) {
        head = newNode;
        phead = newNode;
        if (state == 1)
            head->alloc_count++;
        if (state == 2)
        {
            head->free_count++;
        }
        return;
    } else {
        temp = phead;
        while (temp != NULL) {
            if ((0 == strcmp(temp->name, newNode->name)))
            {
                if (state == 1)
                {
                    temp->alloc_count++;
                }
                if (state == 2)
                {
                    temp->free_count++;
                }
            }
            free(newNode);
            return;
        }
    }
}

```

```

        }
        temp = temp->next;
    }
    temp = phead;
    while (temp->next != NULL) {
        temp = temp->next;
    }
    temp->next = newNode;
    if(state == 1)
    {
        newNode->alloc_count++;
    }
    if(state == 2)
    {
        newNode->free_count++;
    }
}
}

//=====

void insertionSort(Node** head_ref) {
    Node* sorted = NULL;
    Node* current = *head_ref;

    while (current != NULL) {
        Node* next = current->next;
        if (sorted == NULL || sorted->alloc_count >= current->alloc_count) {
            current->next = sorted;
            sorted = current;
        } else {
            Node* temp = sorted;
            while (temp->next != NULL && temp->next->alloc_count < current->alloc_count) {
                temp = temp->next;
            }
            current->next = temp->next;
            temp->next = current;
        }
        current = next;
    }
    *head_ref = sorted;
}

```

```

//=====

void printList(Node* head) {

```

```

unsigned long int total_alloc =0;
unsigned long int total_free =0;

while (phead != NULL) {
//      printf("=====\n");
//      printf("%s\n", phead->name);    // 打印对应 task name
//      printf("allos %lu times\n", phead->alloc_count);    // 打印对应 task alloc 的次数
//      printf("free %lu times\n", phead->free_count);    // 打印对应 task free 的次数
//      total_alloc = total_alloc + phead->alloc_count;    // 计算所有 task 加起来 alloc
//      total_free = total_free + phead->free_count;    // 计算所有 task 加起来 free 的次数
//      if (phead->alloc_count > phead->free_count)
//      {
//          if ( (phead->alloc_count - phead->free_count) > LEAK_TIMES )    // 大量数据的 ftrace
//          log 可筛选泄漏次数大于 LEAK_TIMES 次的 task, 可自行调整打印输出
//          printf("%s leak %lu times\n", phead->name,(phead->alloc_count -
//          phead->free_count));
//      }
//      phead = phead->next;
//  }
//  printf("\n");
//  printf("=====\n");
//  printf("====total alloc %lu=====\n",total_alloc);
//  printf("====total free %lu=====\n",total_free);
//  printf("====total leak %lu=====\n",(total_alloc-total_free));
}

int add_to_list(char * name,int state) {
    if (state==1 || state==2)
    {
        appendNode(name , state);
    }
    return 0;
}

void free_Node()
{
    printList(phead);
    Node* current = head;
    while (current != NULL) {
        Node* next = current->next;
        free(current);
        current = next;
    }
}

```



```

    }
}
int main(int argc, char **argv) {
    char * debug_task_name;
    int state=0;
    FILE *file = fopen(argv[1], "r");
    if (file == NULL) {
        perror("Error opening file");
        return 1;
    }
    char line[1024];
    while (fgets(line, sizeof(line), file)) {
        state = check_alloc_or_free(line);
        debug_task_name = (void *) get_task(line);
        if((debug_task_name != NULL )&& (state ==1 || state ==2))
        {
            add_to_list(debug_task_name,state);
        }
    }
    printf("=====\n");
    insertionSort(&phead);
    printf("=====\n");
    free_Node();
    fclose(file);
    return 0;
}

```